

Name:yamini nakka

[gmail:yaminin90@gmail.com](mailto:yaminin90@gmail.com)

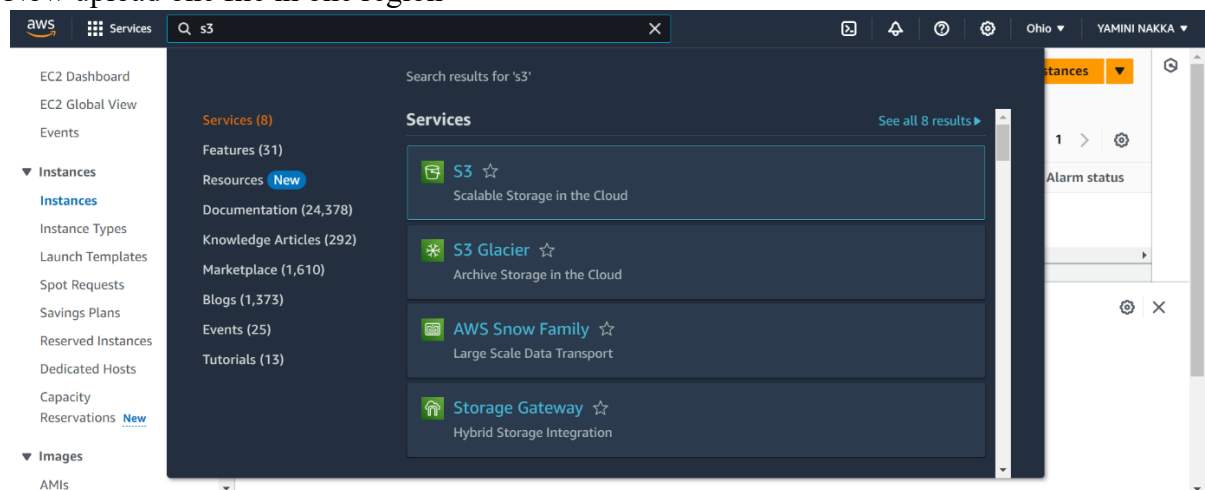
Batch:121

Create two s3 bucket in different regions,upload a file in one s3 bucket and transfer the same file from one s3 bucket to other s3 bucket

S3: S3 is a universal namespace, i.e., the names must be unique globally. Bucket contains a DNS address. Therefore, the bucket must contain a unique name to generate a unique DNS address.

Cross replication: When the cross-region replication is enabled, an object which is uploaded to a specific S3 bucket gets automatically duplicated into a destination bucket that is located at a different AWS location. The objects in the source and destination buckets are copied in an asynchronous manner.

- Create two s3 buckets
- Now upload one file in one region



Pic-1

Step-2

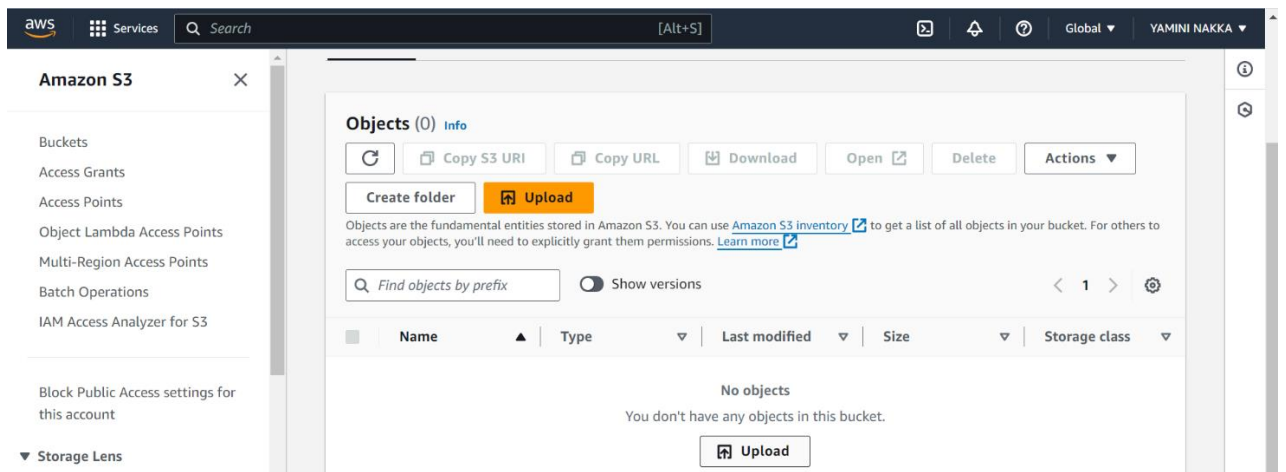
First of create two s3 buckets with different regions

Block Public Access settings for this account	<input type="radio"/>	aws-122	Asia Pacific (Mumbai) ap-south-1	Bucket and objects not public	March 28, 2024, 11:03:03 (UTC+05:30)
Storage Lens	<input type="radio"/>	aws-124	US East (N. Virginia) us-east-1	Bucket and objects not public	March 28, 2024, 11:22:59 (UTC+05:30)
Dashboards					

Pic-2

Step-3

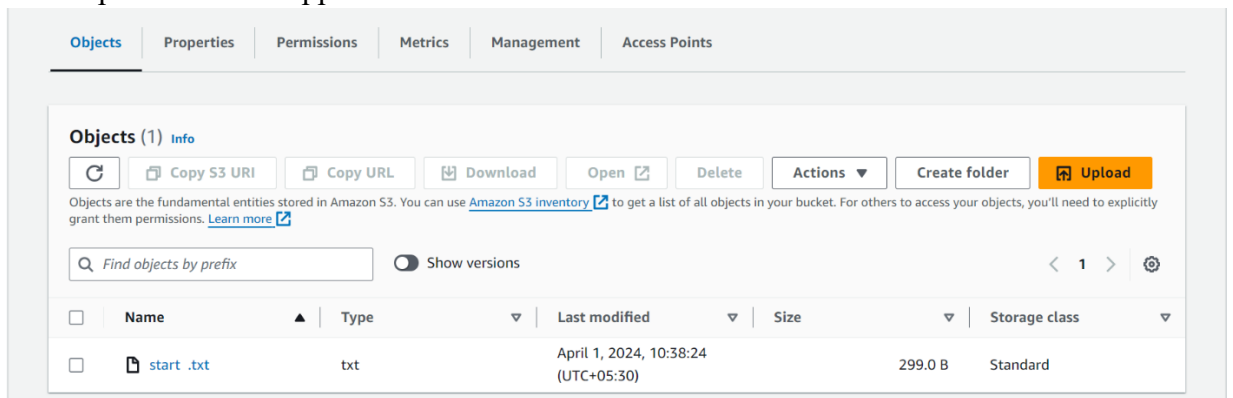
Now you have to upload one file in one region in bucket aws-122



Pic-3

Step-4

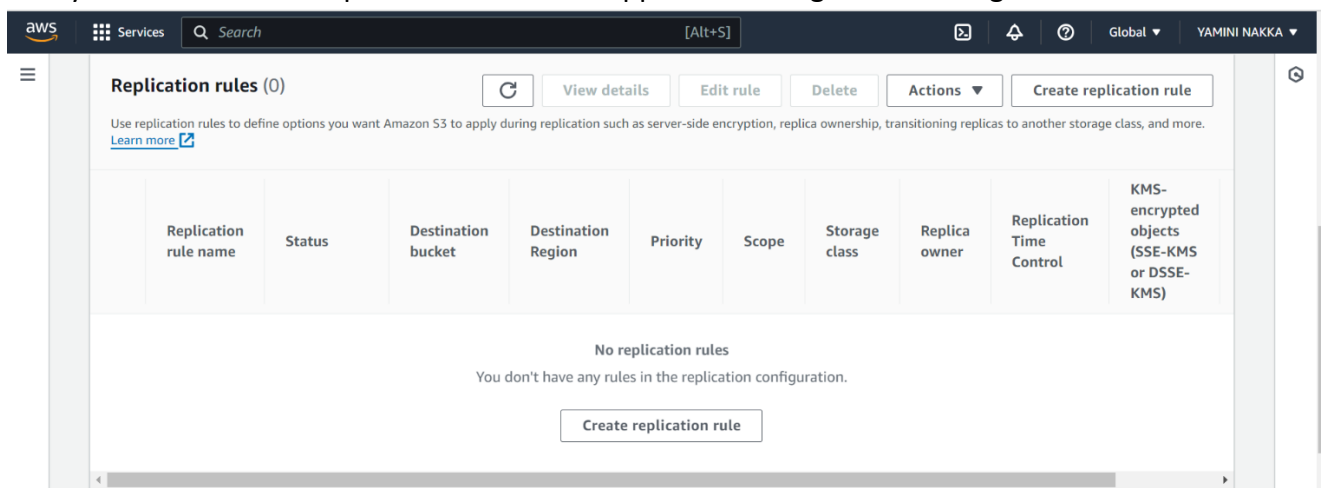
Now the uploaded file is appeared here



Pic-4

Step-5

Now you have to create replication rule .It will appear in management settings



Pic-5

Step-6

Now the replication rule is created

Create replication rule [Info](#)

Replication rule configuration

Replication rule name
rule-1
Up to 255 characters. In order to be able to use CloudWatch metrics to monitor the progress of your replication rule, the replication rule name must only contain English characters.

Status
Choose whether the rule will be enabled or disabled when created.
☒ Enabled
☐ Disabled

Priority
The priority value resolves conflicts that occur when an object is eligible for replication under multiple rules to the same destination. The rule is added to the configuration at the highest priority and the priority can be changed on the replication rules table.
0

Source bucket

Source bucket name
aws-122

Source Region
Asia Pacific (Mumbai) ap-south-1

Choose a rule scope
☐ Limit the scope of this rule using one or more filters
☒ Apply to all objects in the bucket

Destination

Destination
You can replicate objects across buckets in different AWS Regions (Cross-Region Replication) or you can replicate objects across buckets in the same AWS Region (Same-Region Replication). You can also specify a different bucket for each rule in the configuration. [Learn more](#)

[Choose a bucket in this service](#)

Pic-6

IAM role

☒ Create new role
☐ Choose from existing IAM roles
☐ Enter IAM role ARN

Encryption
Server-side encryption protects data at rest.

☐ Replicate objects encrypted with AWS Key Management Service (AWS KMS)
Requires SSE-KMS and SSE-KMS encrypted objects.

Destination storage class
Amazon S3 offers a range of storage classes designed for different use cases. [Learn more](#) or see [Amazon S3 pricing](#)

☐ Change the storage class for the replicated objects

Additional replication options

☒ Replication Time Control (RTC)
Replication Time Control replicates 99.99% of new objects within 15 minutes and includes replication metrics. Additional fees will apply. [Learn more](#)

☒ Replication metrics
With replication metrics, you can monitor the total number and size of objects that are pending replication, and the maximum replication time to the destination region. You can also view and diagnose replication failures. CloudWatch metrics free apply. [Learn more](#) or see [Amazon CloudWatch pricing](#)

☒ To publish event notifications to a destination whenever replication events occur, set S3 event notifications or CloudWatch alarms before replication begins. [Learn more](#)

☐ Delete marker replication
Delete markers created by S3 delete operations will be replicated. Delete markers created by lifecycle rules are not replicated. [Learn more](#)

Pic-7

Step-7

You have to enable the yes,existing the objects

If changes to the configuration aren't displayed, choose the refresh button. Changes apply only to new objects. To replicate existing objects with this configuration, choose **Create replication job**.

Replication rules [Info](#)

Replication enables automatic and asynchronous copy of objects from a source bucket to a destination bucket. Replication can be applied to a group of objects during replication.

Replication configuration settings
Configuration settings affect all replication rules in the bucket.

Source bucket
aws-122

Source Region
Asia Pacific (Mumbai) ap-south-1

Replicate existing objects?

You can enable a one-time Batch Operations job from this replication configuration to replicate objects that already exist in the bucket and to synchronize the source and destination buckets. [Learn more](#) or see [pricing](#)

Existing objects
☐ No, do not replicate existing objects.
☒ Yes, replicate existing objects.

Cancel Submit

Pic-8

Step-8

We have to create job application in that application we have to browse the destination bucket

Job settings

A job is used to execute batch operations on a list of S3 objects. The list of objects is contained in a replication manifest object generated by S3.

Job run options

You can choose whether to have the job start automatically after the replication manifest is generated or to have the job wait in the *Awaiting your confirmation to run* status until you run the job.

☒ Automatically run the job when it's ready
When selected, the job automatically runs without waiting for you to start it.

☐ Wait to run the job when it's ready
Recommended if you want to review the manifest or job details before running the job.

Completion report

Generate a CSV completion report that lists your target objects, task success or error codes, outputs, and descriptions. Completion reports are encrypted using SSE-S3. [Learn more](#)

☒ Generate completion report

Completion report scope

☐ Failed tasks only

☒ All tasks

Completion report destination

Specify a general purpose bucket location to store the completion report. `{job-id}/report.json` will automatically be appended to the specified destination. [Learn more](#)

Format: `s3://bucket[/optional-prefix-with-path-].S3` will append the path with a `"/"`. If you add a `"/"` to the prefix, it will appear as an extra folder in the S3 console.

Permissions

Choose an IAM role with the [required access permissions and trust relationships](#). An IAM role policy template based on your job configuration, and the IAM trust policy required for batch operations to assume the IAM role are available below. [Learn more about IAM roles](#)

☒ View IAM role policy template and IAM trust policy

☒ Create new role

☐ Choose from existing IAM roles

☐ Enter IAM role ARN

Pic-9

Step-9

Here the job application is created

Job ID	Status	Description	Operation	Date created	Total objects	% Complete	Total failed (rate)	Priority
078af782-66ec-42c8-bbec-23fb02107a52	Completed	2024-04-01 - Replicate	Replicate	April 1, 2024, 10:43:53 (UTC+05:30)	4	100%	0 (0%)	10

Pic-10

Step-10

Now the file in one bucket that will appear in another bucket in different region that is cross replication

Amazon S3

Buckets

Access Grants

Access Points

Object Lambda Access Points

Multi-Region Access Points

Batch Operations

IAM Access Analyzer for S3

Block Public Access settings for this account

Storage Lens

Dashboards

Storage Lens groups

AWS Organizations settings

Objects (2) Info

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

☐ Show versions

Name	Type	Last modified	Size	Storage class
job-078af782-66ec-42c8-bbec-23fb02107a52/	Folder	-	-	-
start.txt	txt	April 1, 2024, 10:38:24 (UTC+05:30)	299.0 B	Standard

Pic-11